

1 Cover Letter

2 **Manuscript title - ~~La~~TeXtemplate for Computers & Geosciences**

3 Author 1, Author 2, Author 3

4 Dear Editors-in-Chief,

5

6 please find the enclosed manuscript "..." which we are submitting for exclusive consideration for publication
7 in Computers & Geosciences. We confirm that the submission follows all the requirements and includes all
8 the items of the submission checklist.

9

10 The manuscript presents ...

11

12 We provide the source codes in a public repository with details listed in the section "Code availability".

13

14 Thanks for your consideration.

15

16 Sincerely,

17

18 Authors names

19 Corresponding author affiliation and e-mail

20

21 **Delete before submission:**

22 Please confirm that your submission follows all the requirements of the guidelines, including the submission
23 checklist:

24 - Cover letter

25 - Highlights

26 - Authorship statement

27 - The manuscript must be single column and double spaced

28 - Reference must be in the author-date format

29 - Code availability section

30 *The manuscripts that do meet the requirement guidelines will be desk-rejected.

31 Highlights

32 Manuscript title - L^AT_EXtemplate for Computers & Geosciences

33 Author 1, Author 2, Author 3

- 34 • Phase-1 Hard Negative Mining exploits an intermediate checkpoint to mine visual confounder candi-
35 dates that a fully-converged model cannot provide.
- 36 • River images are the dominant false-positive confounder; per-category FP rates with Clopper-Pearson
37 CIs quantify this precisely.
- 38 • P1-HNM simultaneously improves EfficientNetB0 flood recall from 97.8 % to 99.1 % and reduces river
39 FP rate from 9.2 % to 5.3 %.
- 40 • Recall-first evaluation using PR-AUC alongside per-category FP rates with Clopper-Pearson CIs re-
41 veals per-threshold operating-point differences invisible in aggregate accuracy.
- 42 • A controlled three-way ablation isolates difficulty-ranked selection as the source of improvement over
43 extended training and random injection.

50 ARTICLE INFO

52

53 *Keywords:*

54 Flood detection

55 Hard negative mining

56 Transfer learning

57 Precision-recall

58 Visual confounders

59 Convolutional neural networks

60

61

62

63

64

65

66

67

68

69

70

71

72

73

74

75

76

77

78

79

80

81

82

83

84

85

86

87

88

89

90

91

92

93

94

95

96

97

98

99

100

101

102

103

104

105

106

107

108

109

110

111

112

113

114

115

116

117

118

119

120

121

122

123

124

125

126

127

128

129

130

131

132

133

134

135

136

137

138

139

140

141

142

143

144

145

146

147

148

149

150

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

178

179

180

181

182

183

184

185

186

187

188

189

190

191

192

193

194

195

196

197

198

199

200

201

202

203

204

205

206

207

208

209

210

211

212

213

214

215

216

217

218

219

220

221

222

223

224

225

226

227

228

229

230

231

232

233

234

235

236

237

238

239

240

241

242

243

244

245

246

247

248

249

250

251

252

253

254

255

256

257

258

259

260

261

262

263

264

265

266

267

268

269

270

271

272

273

274

275

276

277

278

279

280

281

282

283

284

285

286

287

288

289

290

291

292

293

294

295

296

297

298

299

300

301

302

303

304

305

306

307

308

309

310

311

312

313

314

315

316

317

318

319

320

321

322

323

324

325

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

378

379

380

381

382

383

384

385

386

387

388

389

390

391

392

393

394

395

396

397

photographs via smartphone applications, social media platforms, or fixed traffic camera feeds, the resulting image stream can provide hyperlocal evidence of active flooding that precedes official event reports by hours (Esparza et al. (2022)). The first-pass screening problem is binary: given an image of unknown provenance and unknown quality, does it depict active flooding? A screening system that operates at high recall on this binary problem enables downstream triage, reducing the burden on human analysts who must validate alerts in real time.

The screening objective is recall-first. A missed flood image (false negative) means that an affected location is not escalated for human review; the consequence is delayed response and potentially preventable harm. A false positive triggers unnecessary resource allocation but carries a lower immediate cost. This asymmetric cost structure makes flood recall—the fraction of actual flood images correctly identified—the primary operational metric. Performance depends strongly on the evaluation metric chosen, particularly given this recall-focused objective. Yet most published flood detection systems report ROC-AUC or overall accuracy, neither of which surfaces per-threshold operating-point differences between architectures nor per-category false positive rates on confounding categories. For a recall-first screening objective, PR-AUC and per-category FP analysis are the appropriate diagnostic tools (Davis and Goadrich (2006); Saito and Rehmsmeier (2015)). We demonstrate this empirically in Section 5.1: despite comparable overall accuracy ($\approx 98.5\%$) and PR-AUC values (0.9976 vs. 0.9991), the two architectures differ by 2 missed floods (9 vs. 7) at $\tau=0.5$ —a per-threshold difference invisible in either aggregate metric but fully captured by per-category false positive analysis.

Visual confounders—rivers, swimming pools, park water features, and wet roads—share water texture with flooded streets and are the dominant source of false positives; per-category analysis is absent from existing benchmarks.

The hard negative mining signal problem. Standard offline hard negative mining (Viola and Jones (2004); Shrivastava et al. (2016)) addresses this by identifying the hardest-scoring non-flood images from a trained model and injecting them into retraining to strengthen the decision boundary. However, this approach fails on fully converged two-phase fine-tuned models for a structural reason: by the end of Phase 2, the model assigns near-zero flood probability to all training non-flood images regardless of their visual similarity to floods. The model has memorised the training distribution; mining from this checkpoint produces empty candidate lists. We address this by mining from the Phase-1 intermediate checkpoint, at which point the model has trained only the classification head and the top backbone layers. At this stage, the early ImageNet feature extractors are still frozen, leaving residual uncertainty on categories that share visual features with the positive class. The Phase-1 checkpoint retains genuine confusion on river and pool

images that the Phase-2 checkpoint has resolved—this is the mining signal. This is a form of *offline* hard negative mining: candidates are identified in a single pre-training pass and injected into the dataset before retraining, rather than re-selected within each mini-batch.

Research gap. State-of-the-art systems (Lee and Joo, 2025; Khan et al., 2023) achieve high aggregate performance but none measure per-category confounder FP rates with confidence intervals, and none ablate the difficulty-ranking contribution of HNM against extended training and random data injection—leaving it unclear whether any improvement is attributable to difficulty ranking or merely additional epochs.

Contributions.

1. **Recall-first evaluation with per-category FP analysis.** We report PR-AUC as the primary metric and provide per-category FP rates with Clopper–Pearson 95 % confidence intervals on 4,099 deduplicated street-level images across eight fine-grained categories, demonstrating that aggregate metrics (accuracy, ROC-AUC) do not surface per-threshold operating-point differences relevant to the recall-first screening objective.
2. **Phase-1 HNM pipeline with controlled ablation.** A three-stage offline pipeline that mines difficulty-ranked confounder candidates from the Phase-1 intermediate checkpoint, applies 5 $\times$ data augmentation, and retrains from the Phase-2 checkpoint—ablated against extended training and random injection to isolate the difficulty-ranking contribution.
3. **Practitioner guidance.** Empirical recommendations for metric selection, architecture choice, and HNM deployment for street-level flood screening.

2. Related Work

2.1. Ground-level flood detection

Esparza et al. (2022) showed that class imbalance directly affects recall in crowdsourced flood reports. Khan et al. (2023) proposed Flood-ResNet50 reporting 96.43 % accuracy on edge hardware. Lee and Joo (2025) introduced AlleyFloodNet, achieving 97.67 % recall with ConvNeXt-Large. Prior work (Papadimos et al., 2023; Alam et al., 2023) benchmarks aggregate performance but none reports per-category false positive rates for water-confounding categories with confidence intervals.

2.2. Aerial and satellite flood detection

Rahneemoonfar et al. (2020) introduced FloodNet, a semantic segmentation dataset from aerial post-Harvey imagery. Tellman et al. (2021) provide a global satellite-based framework estimating flood exposure

from Landsat and MODIS. Street-level screening is complementary: it provides hyperlocal, real-time evidence before satellite revisit cycles update.

2.3. Hard negative mining

Hard negative mining has a long history: Viola and Jones (2004) demonstrated bootstrapped negative mining for face detection, and Shrivastava et al. (2016) systematised online hard example mining (OHEM). Our approach differs in mining *offline* from an *intermediate* checkpoint rather than a fully converged model: convergence eliminates the difficulty signal, but the Phase-1 checkpoint retains genuine uncertainty on visual confounders. Operating at the data level enables our ablation to isolate difficulty ranking from extended training and random injection.

2.4. Progressive and catastrophic-forgetting-aware transfer learning

Simultaneous fine-tuning of all layers risks catastrophic forgetting (Kirkpatrick et al., 2017); Lyu et al. (2025) demonstrated that staged unfreezing improves stability on small datasets. Our two-phase protocol follows this design, and critically, the Phase-1 intermediate checkpoint captures a representation state informative for mining before overfitting the training distribution.

2.5. Evaluation metrics for imbalanced classification

Davis and Goadrich (2006) showed that PR-AUC better reflects recall-oriented performance when the positive class is rare. Saito and Rehmsmeier (2015) empirically confirmed its advantages in class-imbalanced screening regimes. For per-category FP rates on small validation sets we use the exact Clopper–Pearson interval (Clopper and Pearson, 1934), which provides guaranteed coverage regardless of sample size—critical for the swimming pool category ($N=28$, 95 % CI upper bound 12.3 % at zero observed FPs).

3. Dataset

3.1. Source datasets

Flood images are drawn from the USF FloodingDataset University of South Florida (2023), providing binary labels with three severity levels (Major, Moderate, Minor) and non-flood categories including streets, animals, vehicles, buildings, plant life, and parks/walkways. To supplement the dominant confounder, we add river images from the RIWA dataset Wagner et al. (2023), a 1,128-image smartphone/drone/DSLR collection; images are used without segmentation masks, enriching the river category to 399 images (323 train, 76 val).

Table 1

Dataset statistics by category and split. Flood severity labels are Major, Moderate, and Minor. Non-flood categories marked † are water confounders. Note: swimming pool and plant have small validation counts ($N_{\text{val}} < 35$) that limit statistical power for per-category FP rate analysis.

Label	Category	Water?	Train	Val	Total
Flood	street_major	—	640	156	796
Flood	street_moderate	—	245	56	301
Flood	street_minor	—	406	110	516
<i>Flood total</i>			<i>1291</i>	<i>322</i>	<i>1613</i>
Non-flood	river†	Yes	323	76	399
Non-flood	swimming_pool†	Yes	70	28	98
Non-flood	park_walkway†	Partial	326	82	408
Non-flood	street_clear	No	475	105	580
Non-flood	animal	No	233	67	300
Non-flood	building	No	232	55	287
Non-flood	vehicle	No	189	50	239
Non-flood	plant	No	141	34	175
<i>Non-flood total</i>			<i>1989</i>	<i>497</i>	<i>2486</i>
Total			3280	819	4099

3.2. Deduplication

SHA-256 deduplication removed 55 exact-duplicate files from the combined pool, yielding 4,099 unique images. No perceptual deduplication was applied.

3.3. Stratified split

Images are split 80/20 with stratified sampling by fine-grained category label (seed = 42), yielding 3,280 training and 819 validation images. No test split is held out; all reported metrics are on the validation set.

3.4. Dataset statistics

Table 1 reports counts by category and split. Flood prevalence is 39.3 % in both splits (1:1.54 flood:non-flood ratio); class weights are applied during training to compensate.

The swimming pool category ($N_{\text{val}} = 28$) has a Clopper–Pearson 95 % CI upper bound of 12.3 % at zero observed FPs (Section B): a null result cannot be interpreted as confounder robustness.

4. Methods

4.1. Architectures

We evaluate two ImageNet-pretrained convolutional architectures to assess whether Phase-1 HNM is architecture-agnostic. EfficientNetB0 Tan and Le (2019) uses compound scaling (5.3 M parameters); ResNet50 He et al. (2016) uses residual connections (25.6 M parameters). Architecture-specific preprocessing is applied: EfficientNetB0 rescales inputs internally ([0, 255] raw pixels), while ResNet50 uses ImageNet




Figure 1: End-to-end methodology pipeline. Blue boxes: data preparation stages; green boxes: training stages; orange boxes (bold border): Phase-1 HNM novel contribution (Section 4.4); purple box: evaluation. Data flows left-to-right in the top row; the HNM pipeline mines difficulty-ranked candidates from M_1 and retrains from M_2 .

channel-mean subtraction via `resnet50.preprocess_input`.

Both architectures share the same binary classification head:

$$\text{GAP} \cdot \mathbf{x} / \xrightarrow{\text{Drop.0.2/}} \text{Dense.256, ReLU} / \xrightarrow{\text{BN}} \xrightarrow{\text{Drop.0.3/}} \text{Dense.1, } \sigma / \quad (1)$$

where GAP is global average pooling, BN is batch normalisation, and $p \in [0, 1]$ is $P_{\text{non-flood}}$; flood probability = $1 - p$.

4.2. Two-phase progressive fine-tuning

Two-phase progressive fine-tuning mitigates catastrophic forgetting Kirkpatrick et al. (2017); Lyu et al. (2025). Phase 1 trains the classification head and top 30 backbone layers at $\text{LR}_1 = 10^{-4}$ for ≤ 15 epochs (checkpoint M_1). Phase 2 freezes the first 50 backbone layers and fine-tunes all remaining layers at $\text{LR}_2 = 10^{-5}$ for ≤ 20 epochs (checkpoint M_2). Both phases use early stopping (patience = 7 on val loss), `ReduceLROnPlateau`, and balanced class weights. Full hyperparameters are listed in Section A.

4.3. Loss functions

All models are trained with binary cross-entropy (BCE, Eq. 2) combined with balanced class weights to correct for class frequency.

$$\mathcal{L}_{\text{BCE}} = -w_y [y \log p + (1 - y) \log (1 - p)] \quad (2)$$

where w_y is the inverse-frequency class weight for label y .

4.4. Phase-1 hard negative mining

The Phase-1 HNM pipeline proceeds in three stages.

Table 2

Ablation conditions isolating the contributions of additional training and difficulty ranking. All conditions start from the Phase-2 checkpoint M_2 ; they differ only in what data is added to the training set.

Condition	Description	Controls for
HNM (percentile)	Top 20 % hardest negatives by M_1 probability	— (target condition)
Extended training	Same epoch budget; no additional data	Additional training epochs
Random injection	Same count; randomly sampled from flagged categories	Category exposure vs. difficulty ranking

4.4.1. Stage 1: Confounder analysis from the Phase-1 checkpoint

M_1 is run on all `train/non_flood/` images; per-category FP rate is computed at $\tau=0.5$, and any category exceeding $\theta=0.05$ is flagged. Training convergence destroys the mining signal: by Phase 2, EfficientNetB0 assigns flood probability < 0.5 to 99.7 % of training river images (< 2 candidates above threshold across 323 images). The Phase-1 checkpoint has trained only the head and top 30 layers; the frozen early extractors retain genuine uncertainty on visual confounders, providing the difficulty signal Stage 2 exploits.

4.4.2. Stage 2: Difficulty-ranked candidate selection

Images from flagged categories are ranked by descending Phase-1 flood probability p ; the top $\pi=20\%$ are selected. Percentile-based selection is calibration-invariant, ensuring consistent candidate set sizes across architectures. A partition-safety check verifies that no selected hard negative appears in the validation set, preventing data leakage.

4.4.3. Stage 3: Augmented retraining from the Phase-2 checkpoint

Each selected candidate receives 5 $\times$ augmentation (rotation $\pm 30^\circ$, zoom $\pm 30\%$, brightness $[0.6, 1.4]$, horizontal flip) and is injected into the training set. The augmented set is used to retrain from M_2 using the same two-phase protocol, preserving all representation learning from the original training as a warm start.

4.5. Ablation controls

Table 2 defines three conditions to isolate the difficulty-ranking contribution.

If HNM outperforms extended training, additional compute alone cannot explain the gain; if HNM also outperforms random injection, difficulty ranking rather than mere category exposure is the active ingredient.

4.6. Evaluation protocol

PR-AUC is the primary metric (motivation in Section 5.1). At $\tau=0.5$ we also report recall, precision, F1, accuracy, FN and FP counts. Per-category FP rate $r_c = k_c/n_c$ is reported with a Clopper–Pearson 95 %

Table 3

Baseline validation-set performance at $\tau=0.5$ (seed 42; 819 val images: 322 flood, 497 non-flood). Bold marks best per-column value. FN = missed floods; FP = spurious alerts. Flood recall = $TP/(TP+FN)$; Flood Prec. = $TP/(TP+FP)$.

Model	Loss	Acc (%)	Rec (%)	Prec.	F1	FN	FP	PR-AUC	ROC-AUC
EfficientNetB0	BCE	98.3	97.8	0.978	0.978	7	7	0.9976	0.9985
ResNet50	BCE	98.5	97.2	0.991	0.981	9	3	0.9991	0.9995

231 exact binomial CI:

$$[B.\alpha^2; k_c, n_c * k_c + 1 / , B.1 * \alpha^2; k_c + 1, n_c * k_c /] \quad (3)$$

232 where $B.q; a, b /$ is the q -quantile of Beta. $a, b /$ and $\alpha=0.05$. Pairwise comparisons use McNemar's test
 233 McNemar (1947) with Bonferroni correction.

234 5. Experiments and Results

235 5.1. Baseline performance at $\tau = 0.5$

236 Table 3 reports validation performance at $\tau=0.5$ (322 flood, 497 non-flood images). EfficientNetB0
 237 BCE achieves 97.8% flood recall (7 FN, 7 FP; PR-AUC= 0.9976). ResNet50 BCE achieves 97.2% recall
 238 with only 3 FP (PR-AUC= 0.9991, ROC-AUC= 0.9995)— marginally higher discrimination metrics than
 239 EfficientNetB0 but 2 more missed floods at $\tau=0.5$.

240 5.2. Qualitative error analysis

241 Figure 2 shows the worst false positives (red border) and false negatives (orange border) at $\tau=0.5$. Both
 242 architectures' FPs are close-range river scenes visually indistinguishable from flooded roads; FNs are flood
 243 scenes with attenuated water signal (dim lighting, partial flooding in frame). Phase-1 HNM targets these
 244 confounder-driven FP errors directly.

245 5.3. Confounder analysis

246 Table 4 and Figure 3 report per-category FP rates for both BCE baselines. River is the sole significant
 247 confounder: EfficientNetB0 BCE produces 9.2% river FP rate (7/76, CI: [3.8, 17.7] %); ResNet50 BCE
 248 produces 3.9% (3/76, CI: [0.8, 11.0] %). All other categories produce zero FPs at $\tau=0.5$. For swimming
 249 pool ($N=28$), the CI upper bound of 12.3% at zero observed FPs means the null result cannot be interpreted
 250 as confounder robustness.

251 Stage 1 of Phase-1 HNM therefore flags river for EfficientNetB0 ($9.2\% > \theta=5\%$) but not for ResNet50
 252 (3.9%).




Figure 2: Worst false positives (red border, top rows) and false negatives (orange border, bottom rows) at $\tau=0.5$ (val set, seed 42). EfficientNetB0 FPs are visually ambiguous river scenes that share water texture with flood images. ResNet50 FNs include many prototypical flood scenes, indicating systematic rather than boundary-driven failure.




Figure 3: Per-category false positive rate heatmap for both BCE baselines (val set, seed 42, $\tau=0.5$). River is the only category with non-zero FPs in either architecture; EfficientNetB0 BCE river FP rate (9.2%) exceeds ResNet50 BCE (3.9%), identifying it as the primary target for Phase-1 HNM (Table 4).

Table 4

Per-category false positive rates on the validation split, BCE baselines (seed 42, $\tau=0.5$). Clopper–Pearson 95 % binomial CIs via `scipy.stats.beta.interval`. Dagger (†) marks water-confounder categories. Bold upper CI bounds flag categories where a zero observed FP rate should not be interpreted as robustness.

Category	N_{val}	EfficientNetB0 BCE		ResNet50 BCE	
		FP rate	95 % CI	FP rate	95 % CI
river†	76	9.2 %	[3.8, 17.7]	3.9 %	[0.8, 11.0]
plant	34	2.9 %	[0.1, 15.3]	0.0 %	[0.0, 10.3]
swimming_pool†	28	0.0 %	[0.0, 12.3]	0.0 %	[0.0, 12.3]
park_walkway†	82	0.0 %	[0.0, 4.4]	0.0 %	[0.0, 4.4]
animal	67	0.0 %	[0.0, 5.4]	0.0 %	[0.0, 5.4]
building	55	0.0 %	[0.0, 6.5]	0.0 %	[0.0, 6.5]
street_clear	105	0.0 %	[0.0, 3.5]	0.0 %	[0.0, 3.5]
vehicle	50	0.0 %	[0.0, 7.1]	0.0 %	[0.0, 7.1]

Table 5

Phase-1 HNM ablation results (seed 42, $\tau=0.5$). Val Acc from best-epoch checkpoint. River FP rate from per-category confounder analysis. ΔAcc = change vs. respective baseline. Clopper–Pearson 95 % CI in brackets (%). McNemar pairwise tests on river FP status: all $p > 0.05$ after Bonferroni correction.

Model	Condition	Val Acc (%)	ΔAcc	River FP	River FP rate (%)
EfficientNetB0	Baseline	98.29	—	7/76	9.2 [3.8, 17.7]
EfficientNetB0	HNM (percentile)	98.78	+0.49 pp	4/76	5.3 [1.5, 13.0]
EfficientNetB0	Extended training	98.66	+0.37 pp	5/76	6.6 [2.2, 14.9]
EfficientNetB0	Random injection	98.41	+0.12 pp	7/76	9.2 [3.8, 17.7]
ResNet50	Baseline	98.53	—	3/76	3.9 [0.8, 11.0]
ResNet50	HNM (percentile)	<i>pending re-evaluation under corrected preprocessing</i>			

5.4. Phase-1 HNM ablation results

Table 5 reports ablation results; Figure 5 plots river FP rates with Clopper–Pearson CIs; Figure 4 shows PR/ROC curves for all four EfficientNetB0 conditions.

Accuracy is directionally consistent across all conditions: HNM (98.78 %), extended training (98.66 %), random injection (98.41 %), and baseline (98.29 %); this ordering is consistent with hierarchical isolation of the difficulty-ranking contribution, though differences do not reach McNemar significance (see below). Phase-1 HNM simultaneously improves flood recall from 97.8 % to 99.1 % (7 FN $\rightarrow$ 3 FN) and reduces river FP from 9.2 % to 5.3 % (7/76 $\rightarrow$ 4/76)—a joint improvement that neither extended training (6 FN, 5/76 river FP) nor random injection (6 FN, 7/76) achieves. McNemar’s test ($p > 0.05$ after Bonferroni) does not reach significance at $N_{\text{river}}=76$; ≈ 300 river validation images would be needed for 80 % power to confirm the 3–4 pp reduction. ResNet50 HNM ablations are excluded pending re-evaluation under the corrected preprocessing.




Figure 4: PR and ROC curves for the four EfficientNetB0 BCE ablation conditions (val set, seed 42). HNM maintains the highest PR-AUC; accuracy decreases in the order HNM, extended training, random injection, baseline across the full recall range. Curves from the converged Phase-2 checkpoint for each condition.




Figure 5: River false positive rate across all training conditions with Clopper–Pearson 95 % binomial confidence intervals ($N_{river}=76$ in all conditions). The dashed horizontal line marks the Phase-1 HNM flagging threshold $\theta=5\%$. Wide CI overlap confirms that no pairwise difference is statistically significant at the available sample size. Annotated fractions show raw FP counts. Figure generated by `scripts/plot_confounder_fp.py`.

264 6. Discussion

265 6.1. Recall-first evaluation: why per-category FP analysis matters

266 At near-identical accuracy (98.3 % vs. 98.5 %) and similar PR-AUC (0.9976 vs. 0.9991), neither aggregate
267 metric surfaces the 2-FN per-threshold difference between architectures. Per-category FP analysis (Table 4)

provides the operationally relevant diagnostic: EfficientNetB0’s 9.2% river FP rate identifies the specific weakness Phase-1 HNM targets, and the recall difference (7 FN vs. 9 FN) directly quantifies the cost for a recall-first deployment. We recommend all flood screening evaluations report PR-AUC as the primary metric alongside per-category FP rates with Clopper–Pearson CIs, since aggregate accuracy alone does not characterise per-threshold operating-point behaviour on imbalanced categories.

6.2. Why Phase-1 mining works: convergence, uncertainty, and the offline HNM constraint

The Phase-2 checkpoint assigns near-zero flood probability to all training non-flood images, producing empty candidate sets; Phase-1 mining recovers the uncertainty signal before convergence suppresses it. The frozen early extractors retain ImageNet-calibrated features that confuse river and pool images with floods—exactly the candidates needed for boundary refinement. The directionally consistent ordering—HNM, then extended training, then random injection—is consistent with difficulty ranking as the active contribution, though the differences do not reach statistical significance at this validation set size.

6.3. Interpreting the HNM river FP results

Phase-1 HNM reduces river FP from 9.2% to 5.3% ($7/76 \rightarrow 4/76$, a 43% relative reduction) while simultaneously improving flood recall from 97.8% to 99.1% ($7 \text{ FN} \rightarrow 3 \text{ FN}$). This joint improvement is the key result; neither ablation control achieves both objectives. The 3-FP reduction does not reach McNemar significance at $N_{\text{river}}=76$; expanding river validation to ~300 images is needed to confirm the effect.

6.4. Limitations

- **Single seed.** All results are from seed = 42; multi-seed replication is in progress.
- **Small water-confounder N .** Swimming pool ($N_{\text{val}}=28$) and plant ($N_{\text{val}}=34$) are too small for confident confounder conclusions.
- **Geographic bias.** Images are primarily North American; generalisation to tropical, Asian, or African environments is not validated.
- **Phase boundary not ablated.** The (30, 50) freeze boundary is fixed; optimal values for flood-specific data are unknown.
- **No external test set.** All evaluation is on the validation split; out-of-distribution generalisation is not assessed.

6.5. Practitioner recommendations

For teams building street-level flood screening systems:

1. **Report PR-AUC** as the primary metric alongside per-category FP rates with Clopper–Pearson CIs ($N_{\text{val}} < 100$).
2. **Prefer EfficientNetB0 + Phase-1 HNM** at fixed $\tau=0.5$ for recall-first screening (99.1 % recall, 5.3 % river FP); use ResNet50 when threshold calibration is possible.
3. **Apply Phase-1 HNM** when Phase-1 river or pool FP exceeds 5 %; mine from M_1 , not from the converged Phase-2 model.
4. **Use BCE + balanced class weights** as the default loss.
5. **Expand water-confounder validation sets** before deployment, especially swimming pools and lakes.

7. Conclusion

We presented Phase-1 Hard Negative Mining (Phase-1 HNM), a three-stage pipeline exploiting an intermediate training checkpoint as a mining oracle—the core insight being that convergence destroys the mining signal, while the Phase-1 checkpoint retains genuine uncertainty on visual confounders.

On a 4,099-image benchmark, EfficientNetB0 (PR-AUC= 0.9976, 97.8 % recall) and ResNet50 (PR-AUC= 0.9991, 97.2 % recall) achieve strong baselines with complementary per-threshold tradeoffs invisible in aggregate accuracy or PR-AUC; per-category FP analysis identifies river as the sole significant confounder (9.2 % vs. 3.9 %) and the specific target for Phase-1 HNM.

Phase-1 HNM simultaneously improves EfficientNetB0 flood recall from 97.8 % to 99.1 % (7 FN→3 FN) and reduces river FP from 9.2 % to 5.3 %, outperforming extended training and random injection. The river FP reduction is directionally meaningful but does not reach McNemar significance at $N_{\text{river}}=76$; ù300 river validation images are needed to confirm it.

We recommend evaluating flood screening at the per-category FP level with Clopper–Pearson CIs, reporting PR-AUC as the primary metric, and applying Phase-1 mining when any water-confounder exceeds a 5 % training FP rate. Future work should conduct multi-seed replication, expand water-confounder validation sets, and evaluate out-of-distribution generalisation on geographically diverse datasets.

Acknowledgements

[ANONYMISED FOR REVIEW]

325 Data and code availability

326 Training dataset: zinnia82/flood-binary-hnm-benchmark (Hugging Face Hub). River augmenta-
 327 tion images: RIWA dataset Wagner et al. (2023), [https://www.kaggle.com/datasets/franzwagner/](https://www.kaggle.com/datasets/franzwagner/river-water-segmentation-dataset)
 328 river-water-segmentation-dataset. Source code: [TODO — GitHub link before submission].

329 References

- 330 Alam, F., Alam, T., Hasan, M.A., Hasnat, A., Imran, M., Ofii, F., 2023. MEDIC: A multi-task learning dataset for disaster
 331 image classification. *Neural Computing and Applications* 35, 2609–2632. doi:10.1007/s00521-022-07717-0.
- 332 Clopper, C.J., Pearson, E.S., 1934. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*
 333 26, 404–413. doi:10.1093/biomet/26.4.404.
- 334 Davis, J., Goadrich, M., 2006. The relationship between precision-recall and ROC curves, in: *Proc. ICML 2006*, pp. 233–240.
 335 doi:10.1145/1143844.1143874.
- 336 Esparza, M., Farahmand, H., Brody, S., Mostafavi, A., 2022. Examining data imbalance in crowdsourced reports for improving
 337 flash flood situational awareness. URL: <https://arxiv.org/abs/2207.05797>, arXiv:2207.05797.
- 338 He, K., Zhang, X., Ren, S., Sun, J., 2016. Deep residual learning for image recognition, in: *Proc. CVPR 2016*, pp. 770–778.
 339 doi:10.1109/CVPR.2016.90.
- 340 Khan, M.A., Ahmed, N., Padela, J., Raza, M.S., Gangopadhyay, A., Wang, J., Foulds, J., Busart, C.E., Erbacher, R., 2023.
 341 Flood-ResNet50: Optimized deep learning model for efficient flood detection on edge device, in: *Proc. ICMLA 2023*, pp.
 342 506–511. doi:10.1109/ICMLA58977.2023.00077.
- 343 Kirkpatrick, J., Pascanu, R., Rabinowitz, N., Veness, J., Desjardins, G., Rusu, A.A., Milan, K., Quan, J., Ramalho, T.,
 344 Grabska-Barwinska, A., Hassabis, D., Clopath, C., Kumaran, D., Hadsell, R., 2017. Overcoming catastrophic forgetting in
 345 neural networks. *Proceedings of the National Academy of Sciences* 114, 3521–3526. doi:10.1073/pnas.1611835114.
- 346 Lee, O., Joo, H., 2025. AlleyFloodNet: A ground-level image dataset for rapid flood detection in economically and flood-
 347 vulnerable areas. *Electronics* 14, 2082. doi:10.3390/electronics14102082.
- 348 Lyu, S., Lyu, R., Zhao, Y., Gao, W., 2025. ResNet50-based progressive transfer learning, in: *Proc. ICCNSE 2025*, pp. 1–6.
 349 doi:10.1109/ICCNSE66404.2025.11144108.
- 350 McNemar, Q., 1947. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*
 351 12, 153–157. doi:10.1007/BF02295996.
- 352 Papadimos, T., Andreadis, S., Gialampoukidis, I., Vrochidis, S., Kompatsiaris, I., 2023. Flood-related multimedia benchmark
 353 evaluation: Challenges, results and a novel GNN approach. *Applied Sciences* 13, 4814. doi:10.3390/app13084814.
- 354 Rahnemounfar, M., Chowdhury, T., Sarkar, A., Varshney, D., Yari, M., Murphy, R., 2020. FloodNet: A high resolution aerial
 355 imagery dataset for post flood scene understanding. URL: <https://arxiv.org/abs/2012.02951>, arXiv:2012.02951.
- 356 Saito, T., Rehmsmeier, M., 2015. The precision-recall plot is more informative than the ROC plot when evaluating binary
 357 classifiers on imbalanced datasets. *PLOS ONE* 10, e0118432. doi:10.1371/journal.pone.0118432.
- 358 Shrivastava, A., Gupta, A., Girshick, R., 2016. Training region-based object detectors with online hard example mining, in:
 359 *Proc. CVPR 2016*, pp. 761–769. URL: <https://arxiv.org/abs/1604.03540>.
- 360 Tan, M., Le, Q.V., 2019. EfficientNet: Rethinking model scaling for convolutional neural networks, in: *Proc. ICML 2019*, pp.

Table 6

Training hyperparameters verified against source code.

Hyperparameter	Value
Phase-1 learning rate	$1 \dot{\sim} 10^{*4}$
Phase-2 learning rate	$1 \dot{\sim} 10^{*5}$
Phase-1 max epochs	15
Phase-2 max epochs	20
Early stopping patience	7 epochs (val loss)
LR reduce patience (Phase 1)	3 epochs
LR reduce patience (Phase 2)	4 epochs
Batch size	32
Input image size	224 $\dot{\sim}$ 224
Dropout rates	0.2 (pre-Dense), 0.3 (post-BN)
Dense hidden units	256
Phase-1 trainable layers (top n)	30
Phase-2 frozen layers (first n)	50
HNM mining checkpoint	Phase-1 best val loss
HNM retraining start	Phase-2 checkpoint
HNM Phase-1 learning rate	$5 \dot{\sim} 10^{*5}$
HNM Phase-2 learning rate	$1 \dot{\sim} 10^{*5}$
HNM top percentile	20 %
HNM FP flagging threshold θ	0.05
HNM augmentation factor	5 $\dot{\sim}$ per mined image
HNM augmentation operations	Rotation ,30, zoom ,30 $\sim$, brightness [0.6, 1.4], H-flip
Decision threshold τ	0.5
Random seed	42 (single seed; multi-seed pending)
Compute platform	Apple M4 (16 GB, Metal GPU)

361 6105–6114. URL: <https://proceedings.mlr.press/v97/tan19a.html>.

362 Tellman, B., Sullivan, J.A., Kuhn, C., Kettner, A.J., Doyle, C.S., Brakenridge, G.R., Erickson, T.A., Slayback, D.A.,
363 2021. Satellite imaging reveals increased proportion of population exposed to floods. *Nature* 596, 80–86. doi:10.1038/
364 s41586-021-03695-w.

365 University of South Florida, 2023. FloodingDataset: Street-level flood binary classification dataset. Google Drive dataset.
366 URL: <https://drive.google.com/drive/folders/1PXc9VTeQgV5WeNxa2N0C9kbRpnQtg20o>. [CITATION NEEDED — no peer-
367 reviewed paper identified as of April 2026; see TODO_CITATIONS.md].

368 Viola, P., Jones, M., 2004. Robust real-time face detection. *International Journal of Computer Vision* 57, 137–154. doi:10.
369 1023/B:VISI.0000013087.49260.fb.

370 Wagner, F., Eltner, A., Maas, H.G., 2023. River water segmentation in surveillance camera images: A comparative study of
371 offline and online augmentation using 32 CNNs. *International Journal of Applied Earth Observation and Geoinformation*
372 119, 103305. doi:10.1016/j.jag.2023.103305.

373 A. Hyperparameters

374 Table 6 lists all training hyperparameters.

Table 7

Upper bound of 95 % Clopper–Pearson CI for $k=0$ FP observed, by sample size n . Categories in the present study are listed in parentheses. The upper CI bound for swimming pool (12.3 %) implies that a true FP rate of up to 12.3 % is consistent with the observed zero-FP result.

n (val images)	Upper 95 % CI bound
10	30.8 %
15	21.8 %
28	12.3 % (swimming pool)
34	10.3 % (plant)
50	7.1 % (vehicle)
55	6.5 % (building)
67	5.4 % (animal)
76	4.7 % (river)
82	4.4 % (park_walkway)
105	3.5 % (street_clear)

B. Clopper–Pearson CI reference

Table 7 lists the 95 % Clopper–Pearson CI upper bound for zero observed FPs, by sample size.